

Homework 5 实验报告

实验内容：

- 测试sklearn中以下聚类算法在tweets数据集上的聚类效果。
- 使用NMI(Normalized Mutual Information)作为评价指标。

Method name	Parameters	Scalability	Usecase	Geometry (metric used)
K-Means	number of clusters	Very large <code>n_samples</code> , medium <code>n_clusters</code> with MiniBatch code	General-purpose, even cluster size, flat geometry, not too many clusters	Distances between points
Affinity propagation	damping, sample preference	Not scalable with <code>n_samples</code>	Many clusters, uneven cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
Mean-shift	bandwidth	Not scalable with <code>n_samples</code>	Many clusters, uneven cluster size, non-flat geometry	Distances between points
Spectral clustering	number of clusters	Medium <code>n_samples</code> , small <code>n_clusters</code>	Few clusters, even cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
Ward hierarchical clustering	number of clusters	Large <code>n_samples</code> and <code>n_clusters</code>	Many clusters, possibly connectivity constraints	Distances between points
Agglomerative clustering	number of clusters, linkage type, distance	Large <code>n_samples</code> and <code>n_clusters</code>	Many clusters, possibly connectivity constraints, non Euclidean distances	Any pairwise distance
DBSCAN	neighborhood size	Very large <code>n_samples</code> , medium <code>n_clusters</code>	Non-flat geometry, uneven cluster sizes	Distances between nearest points
Gaussian mixtures	many	Not scalable	Flat geometry, good for density estimation	Mahalanobis distances to centers

实验步骤：

1. 读入 tweets

```
def read(file_name):
    X = []
    Y = []
    with open(file_name, 'r', errors='ignore') as f:
        for line in f:
            obj = json.loads(line.strip())
            X.append(obj['text'])
            Y.append(obj['cluster'])
    return X, Y
```

2. 将文本表示为向量

用 sklearn 里 feature_extraction 模块里的 TfidfVectorizer 将文本转换为 tfidf 的向量

```
def feature_extract(X):
    from sklearn.feature_extraction.text import TfidfVectorizer
    vectorizer = TfidfVectorizer()
    return vectorizer.fit_transform(X).toarray()
```

3. 测试 sklearn 的各种聚类效果

用 nmi 评估:

```
def nmi_eval(Y, Y_):  
    from sklearn.metrics import normalized_mutual_info_score  
    return normalized_mutual_info_score(Y, Y_, average_method='arithmetic')
```

测试 k-means:

```
def test_KMeans(X, Y):  
    from sklearn.cluster import KMeans  
    Y_ = KMeans(len(set(Y))).fit_predict(X)  
    print('k-means :', nmi_eval(Y, Y_))
```

结果:

```
k-means : 0.7976727139528518
```

测试 affinity propagation:

```
def test_AffinityPropagation(X, Y):  
    from sklearn.cluster import AffinityPropagation  
    af = AffinityPropagation().fit(X)  
    Y_ = af.labels_  
    print('affinity propagation :', nmi_eval(Y, Y_))
```

结果:

```
affinity propagation : 0.7715980389472071
```

测试 mean shift:

```
def test_MeanShift(X, Y):  
    from sklearn.cluster import MeanShift, estimate_bandwidth  
    bandwidth = estimate_bandwidth(X)  
    ms = MeanShift(bandwidth=bandwidth, bin_seeding=True)  
    ms.fit(X)  
    Y_ = ms.labels_  
    # print(Y_)  
    print('mean shift :', nmi_eval(Y, Y_))
```

结果:

```
mean shift : -8.557811214631876e-17
```

测试 spectral clustering:

```
def test_SpectralClustering(X, Y):  
    from sklearn.cluster import SpectralClustering  
    sc = SpectralClustering(n_clusters=len(set(Y)))  
    sc.fit(X)  
    Y_ = sc.labels_  
    print('spectral clustering :', nmi_eval(Y, Y_))
```

结果:

```
spectral clustering : 0.6828491203018877
```

测试 ward hierarchical clustering:

```
def test_WardHierarchicalClustering(X,Y):  
    from sklearn.cluster import AgglomerativeClustering  
    ward = AgglomerativeClustering(n_clusters=len(set(Y)), linkage='ward')  
    ward.fit(X)  
    Y_ = ward.labels_  
    print('ward hierarchical clustering :', nmi_eval(Y,Y_))
```

结果:

```
ward hierarchical clustering : 0.7797268365583422
```

测试 agglomerative clustering:

```
def test_AgglomerativeClustering(X,Y):  
    from sklearn.cluster import AgglomerativeClustering  
    model = AgglomerativeClustering(n_clusters=len(set(Y)))  
    model.fit(X)  
    Y_ = model.labels_  
    print('agglomerative clustering :', nmi_eval(Y,Y_))
```

结果:

```
agglomerative clustering : 0.7797268365583422
```

测试 DBSCAN:

```
def test_DBSCAN(X,Y):  
    from sklearn.cluster import DBSCAN  
    db = DBSCAN().fit(X)  
    Y_ = db.labels_  
    print(Y_)  
    print('DBSCAN :', nmi_eval(Y,Y_))
```

结果:

```
DBSCAN : 0.036037108078898594
```

测试高斯混合:

```
def test_GaussianMixture(X,Y):  
    from sklearn import mixture  
    clf = mixture.GaussianMixture(n_components=len(set(Y)), covariance_type='full')  
    clf.fit(X)  
    Y_ = clf.predict(X)  
    print('gaussian mixture :', nmi_eval(Y,Y_))
```

运行报错, 提示 MemoryError

只跑了前 500 个数据, 结果:

```
gaussian mixture : 0.8126465671419534
```